Twitter Hate Speech Detection

Can Content Moderation be Automated?

Flatiron School Capstone
Sidney Kung

The Problem of Human Content Moderation

- **Every major tech company** uses third-party contractors
- **Automating** this process could **reduce labor exploitation**
- What is **Hate Speech?**

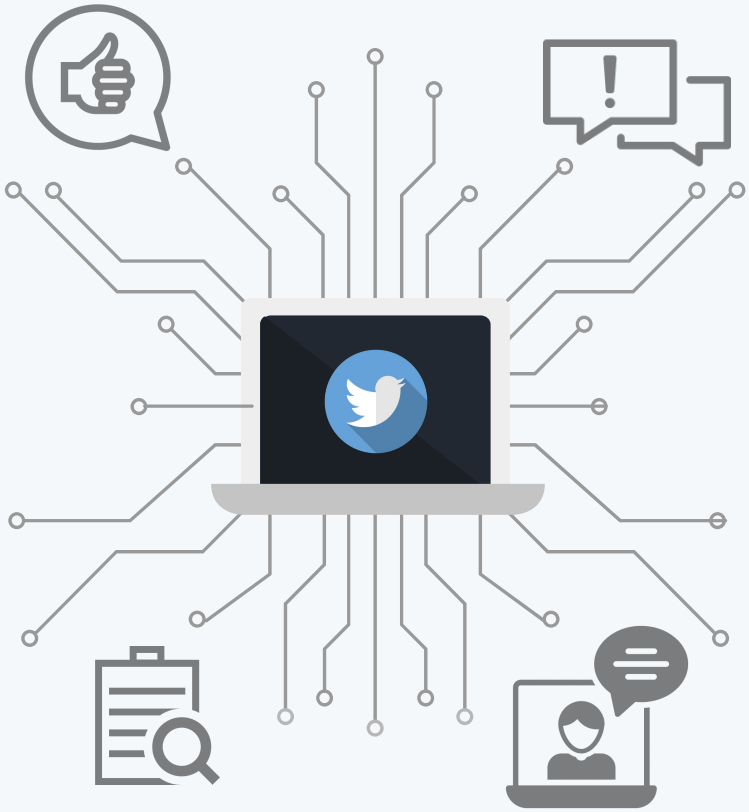

CRISP-DM Process

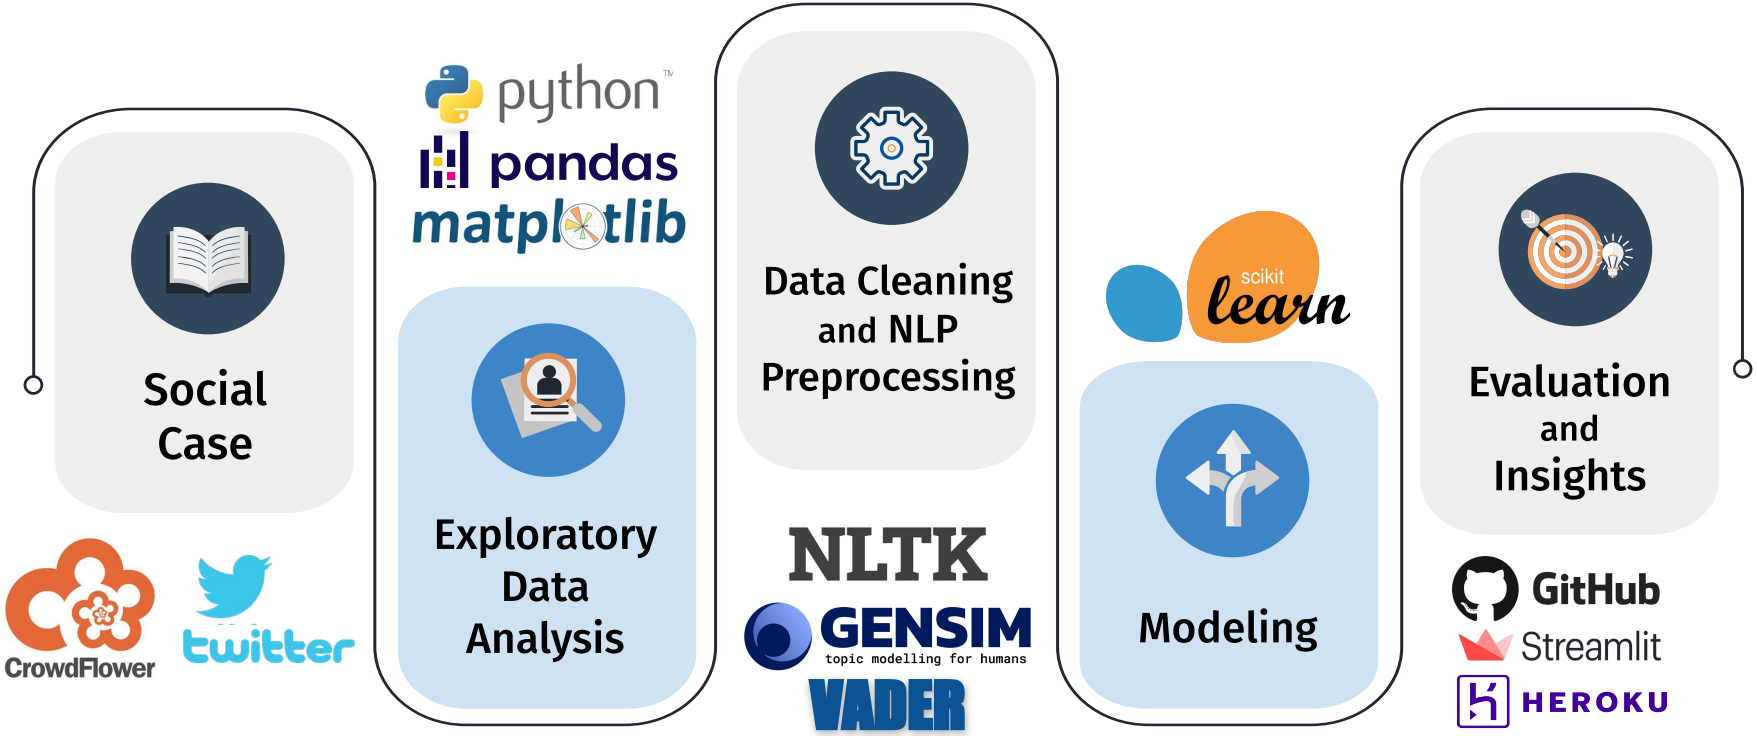

Unique Words per Label

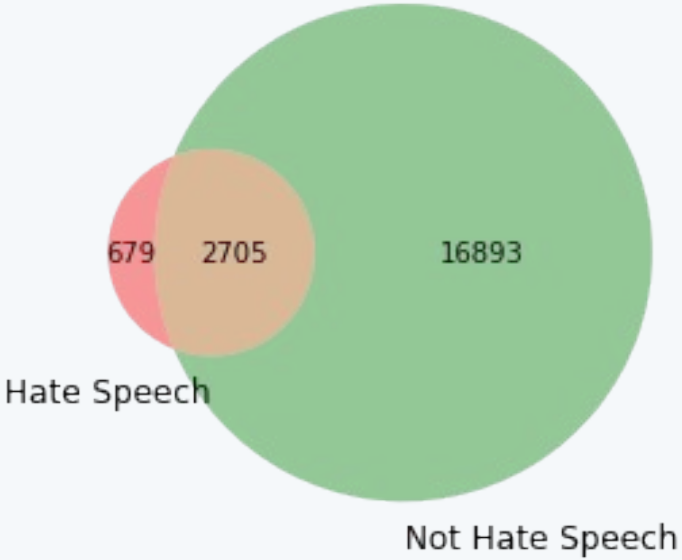

Data Understanding

Sourced from 2017 Cornell University **research study**.

24,802 Tweets

20,277 Word Vocabulary

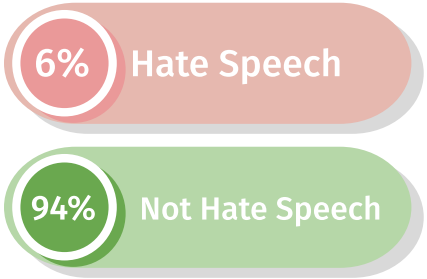

Data Analysis

1. What are the **linguistic differences** between hate speech and offensive language?
2. What are the **popular hashtags** of each tweet type?
3. What is the **overall polarity** of the tweets?

What are the **linguistic differences** between hate speech and offensive language?

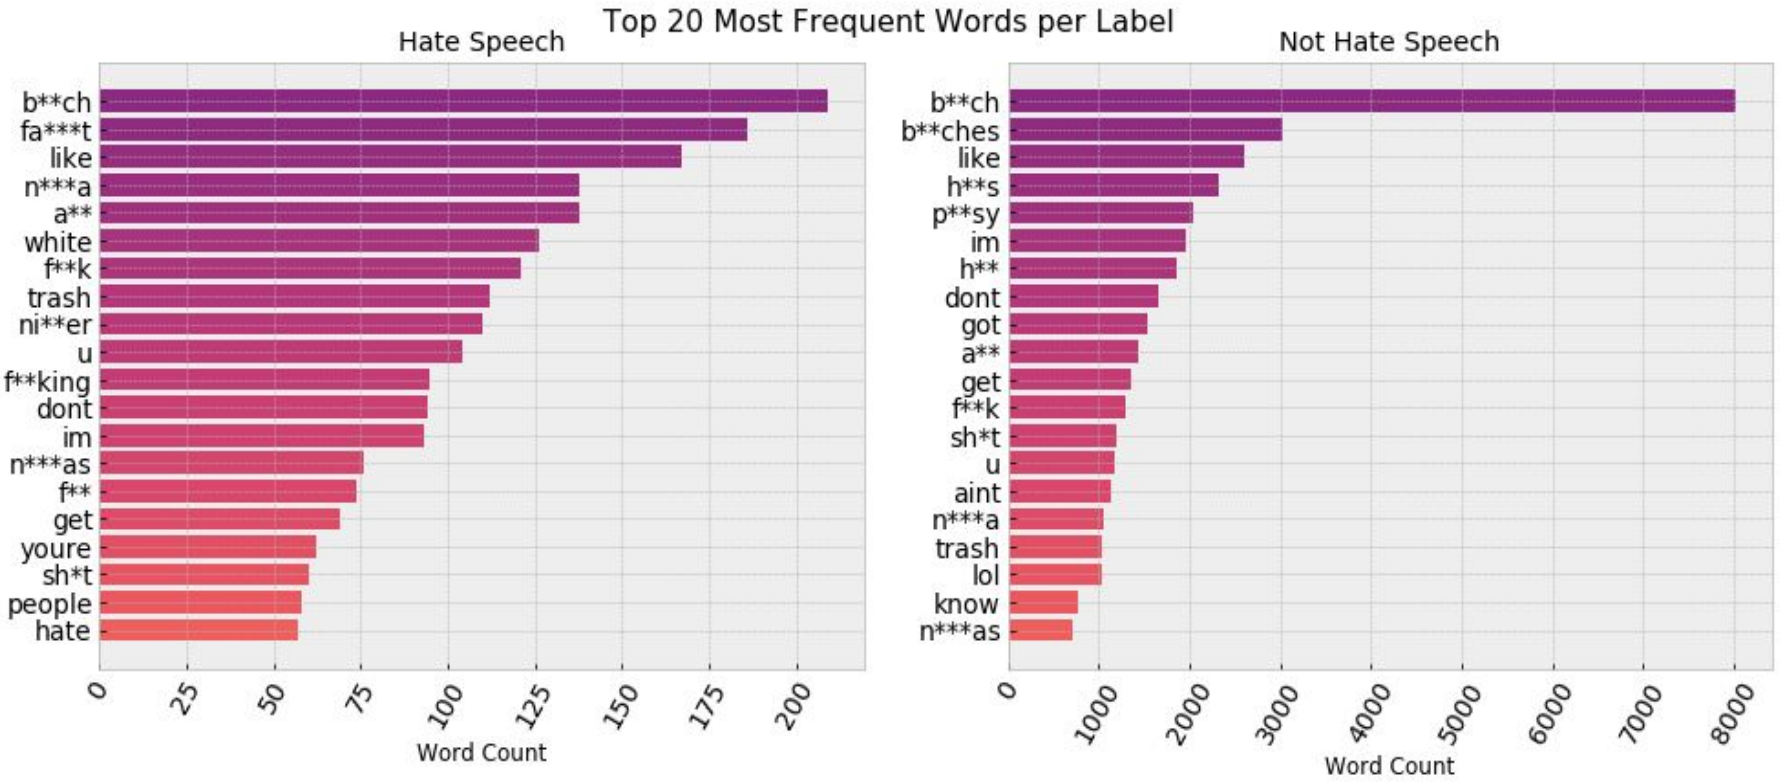

What are the **popular hashtags** of each tweet type?

What is the **overall polarity** of the tweets?

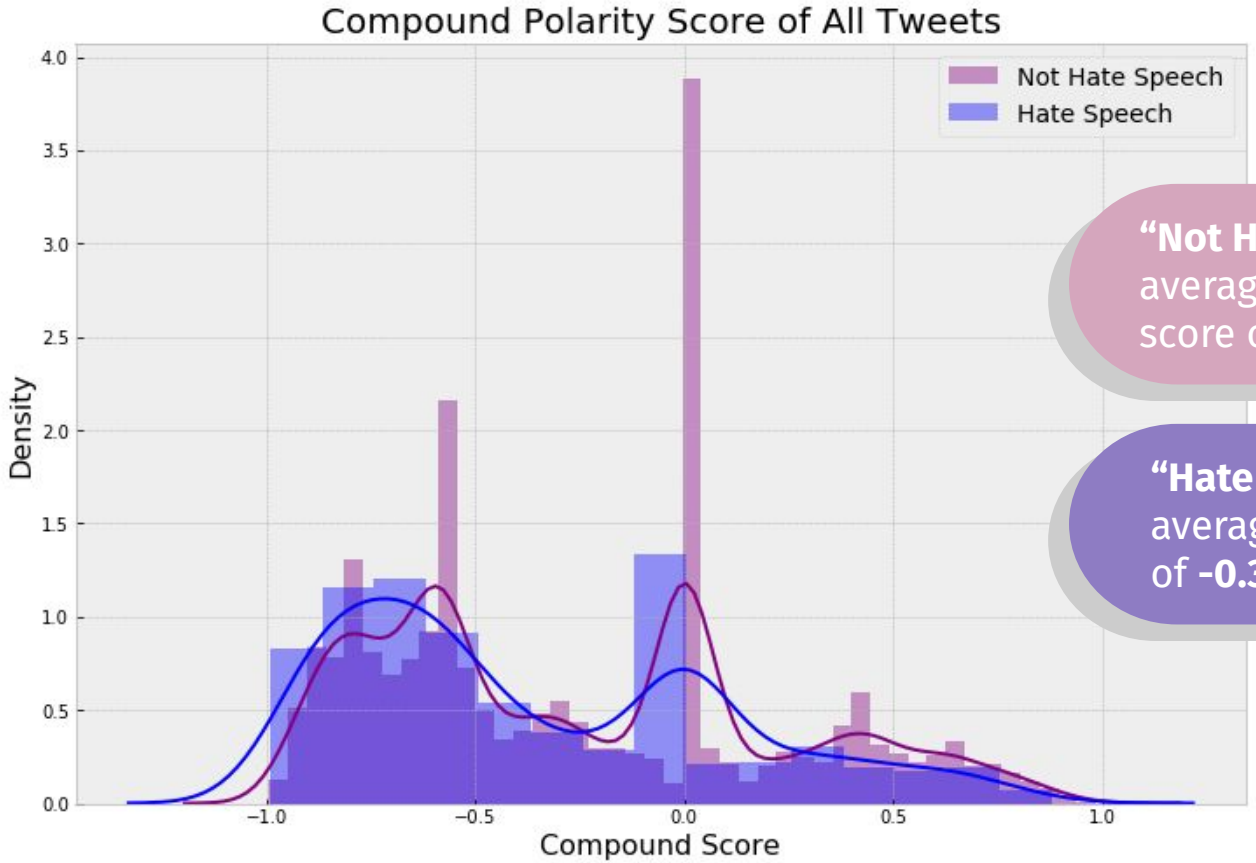

"Not Hate Speech" tweets: average compound score of **-0.263**

"Hate Speech" tweets: average compound score of **-0.363**

Modeling Process

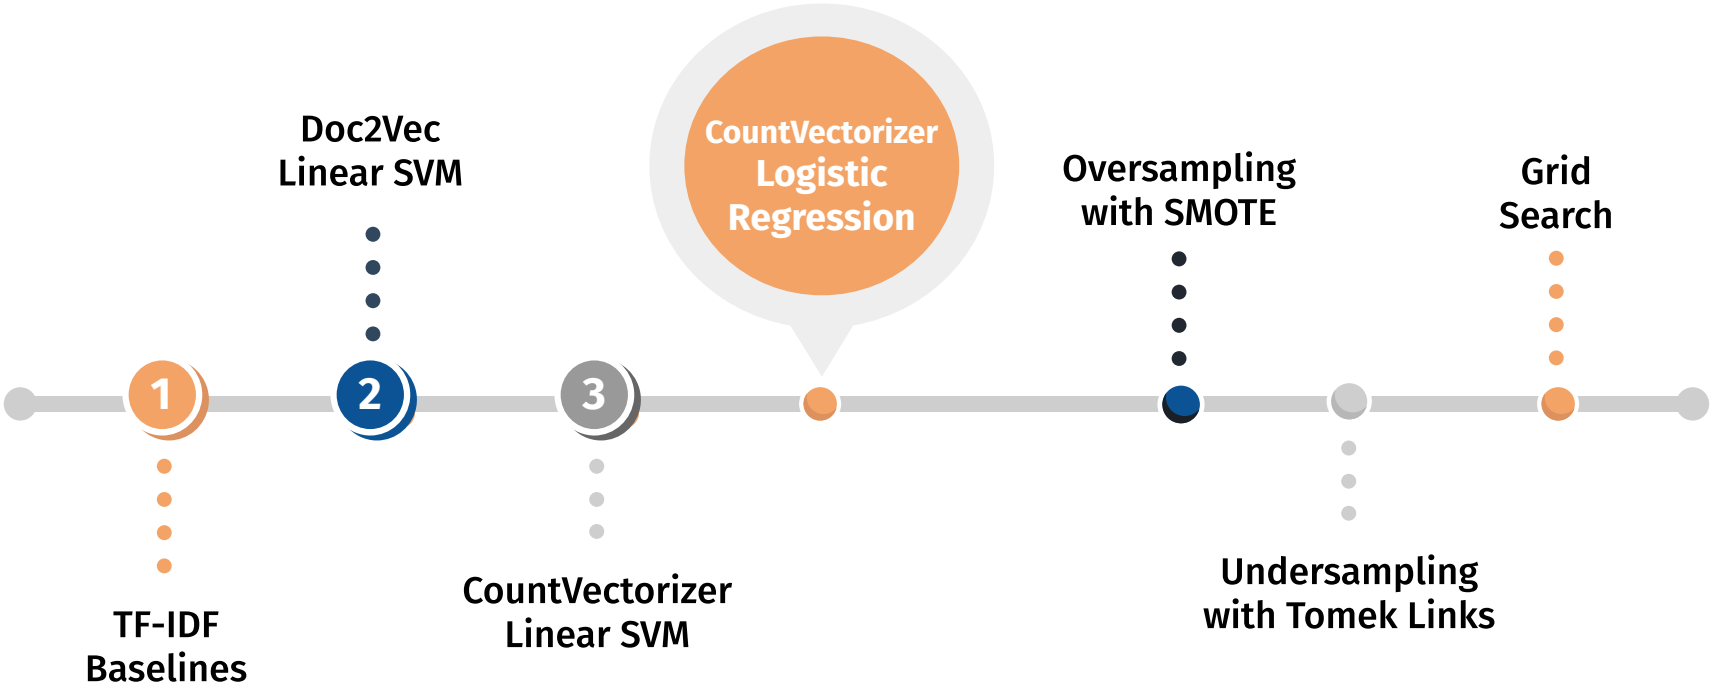

Baseline Model

Naive Bayes
with TF-IDF Vectorization

F1 Score	0.1923
Recall	0.1254

Final Model

Logistic Regression
with Count Vectorization

F1 Score	0.3958
Recall	0.624

Model Deployment

hate-speech-predictor.herokuapp.com

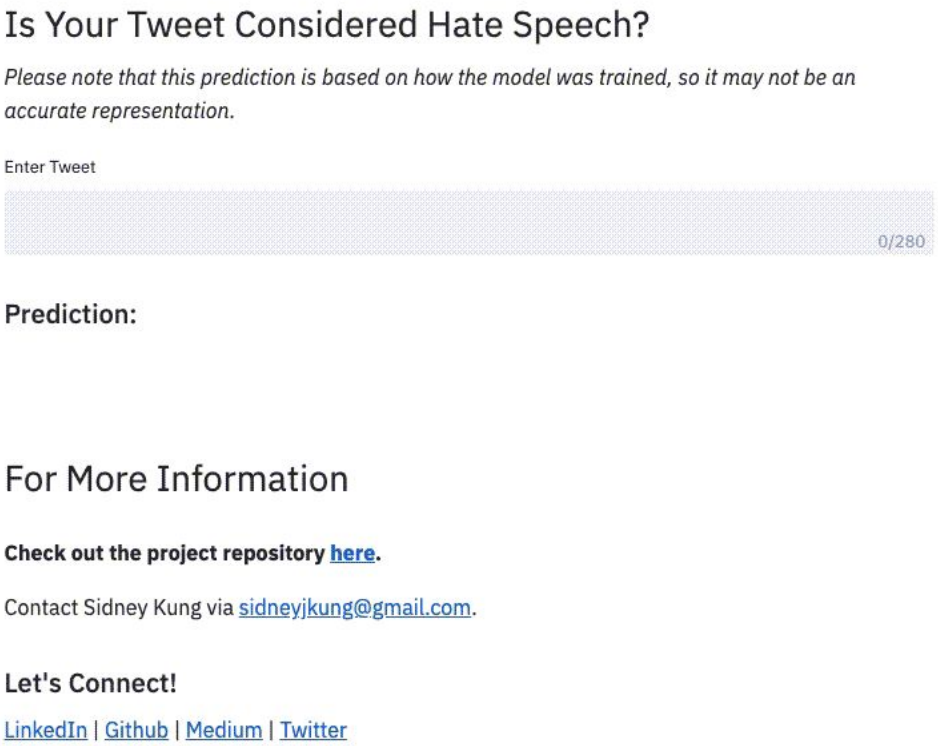

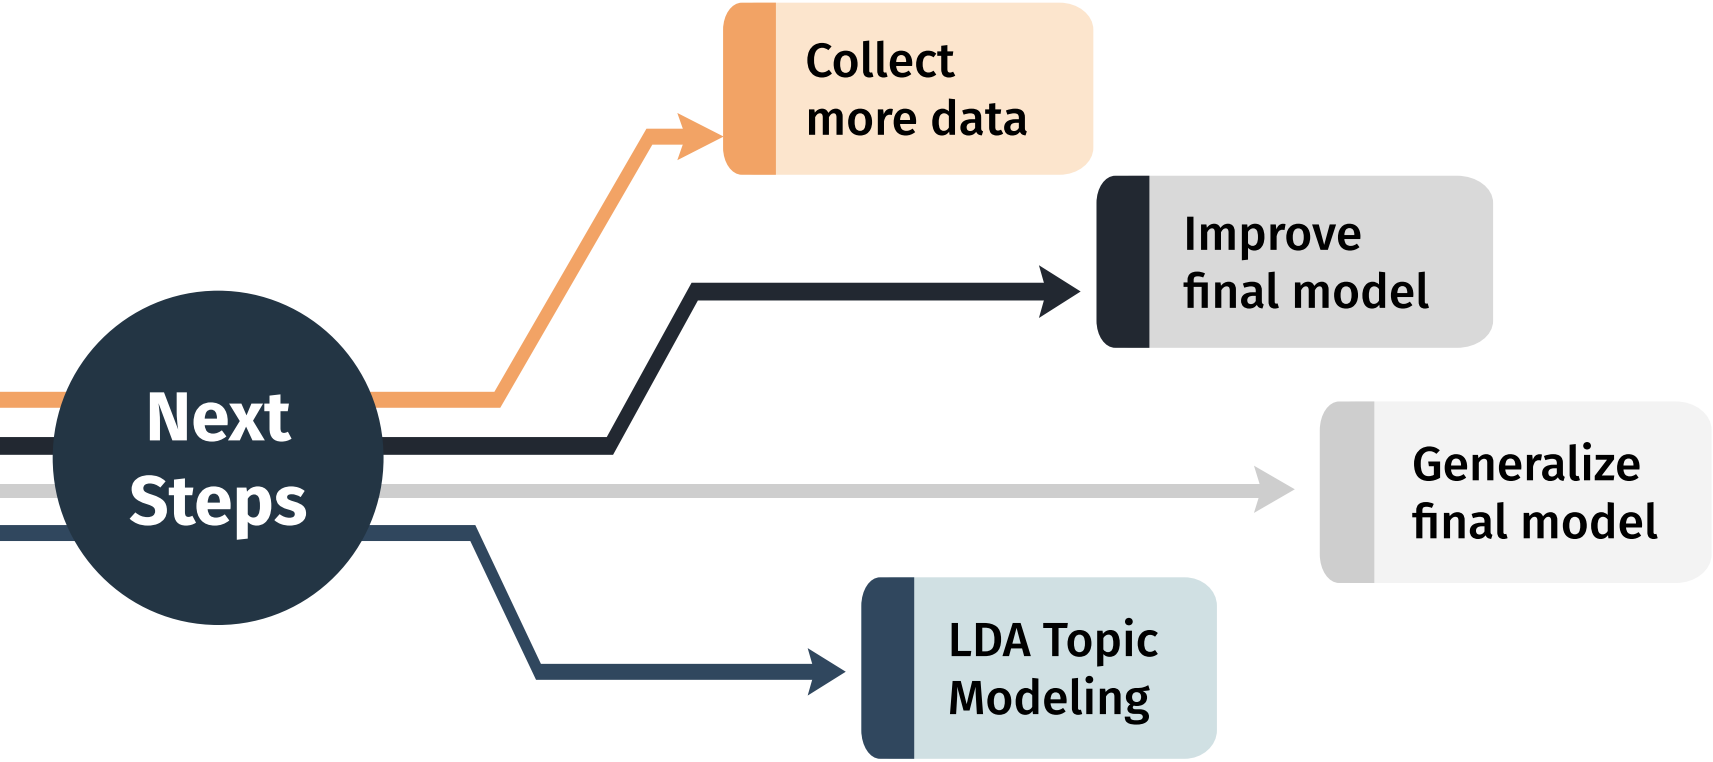
Next Steps
Collect more data
Improve final model
Generalize final model
LDA Topic Modeling

Thank You!

GitHub Repository
github.com/sidneykung/
twitter_hate_speech_detection

Web App on Heroku
hate-speech-predictor.
herokuapp.com

SidneyKung.com

SidneyJKung@gmail.com

linkedin.com/in/sidneykung

@Sidney_K98

Presentation Template:
SlidesGo